

Open in app ↗



★ Member-only story

Understanding Apache Spark Architecture and PySpark Job Execution

8 min read · May 6, 2025



DataWithSantosh

Follow



Listen



Share



More

What is Apache Spark?

Apache Spark is an open-source, distributed computing framework designed for **big data processing**. It's fast, scalable, and handles large datasets across multiple computers (a cluster). Think of Spark as a super-efficient chef who can cook massive meals by coordinating many kitchen assistants (computers) to work together.

Spark is used for:

- **Batch processing** (e.g., processing historical sales data).
- **Stream processing** (e.g., analyzing live social media feeds).
- **Machine learning** (e.g., training models on large datasets).
- **SQL queries** (e.g., querying big data like a database).

Spark's key advantage is its **in-memory processing**, which makes it much faster than older systems like Hadoop MapReduce, as it keeps data in memory rather than constantly reading from disk.




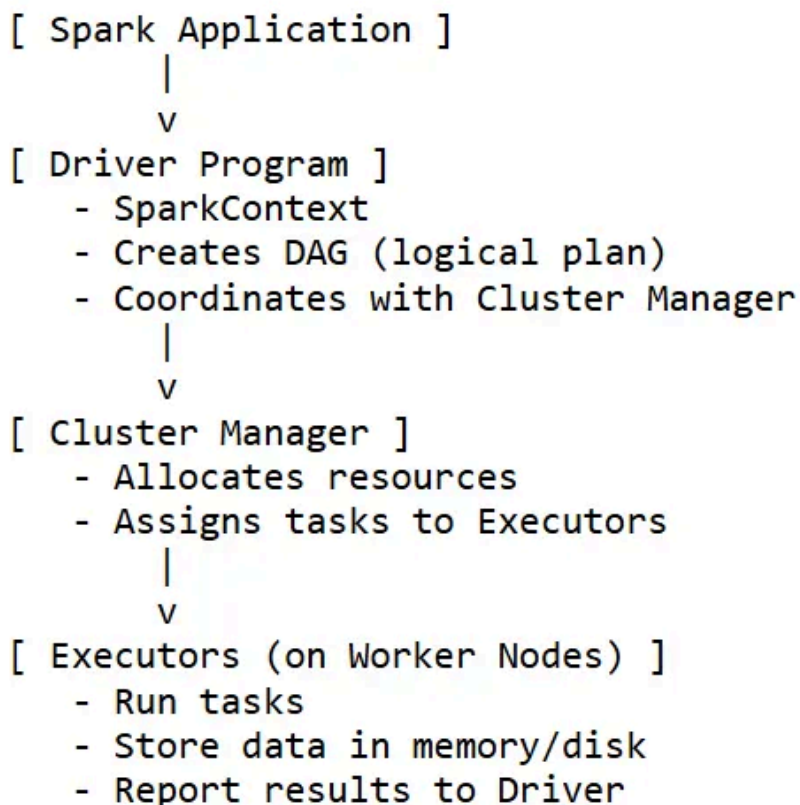
image Source: <https://spark.apache.org/>

Apache Spark Architecture: The Big Picture

Spark's architecture is like a well-organized factory:

- There's a **manager** (Spark Driver) who plans the work.
- A **supervisor** (Cluster Manager) assigns tasks to workers.
- **Workers** (Executors) do the actual work on different machines.

Here's a visual representation of Spark's architecture:



Let's break down each component and how they work together.

Components of Apache Spark Architecture

1. Spark Driver

- **What is it?** The Driver is the brain of a Spark application. It runs the main program and coordinates the entire job.
- **Key Responsibilities:**
- Creates the **SparkContext**, which is like the entry point to Spark. It connects your program to the Spark cluster.
- Converts your code (e.g., Python, Scala) into a **Directed Acyclic Graph (DAG)**, a logical plan of tasks.
- Splits the DAG into **stages** and **tasks** and sends tasks to Executors.
- Tracks the progress of tasks and handles failures.
- **Analogy:** Think of the Driver as a project manager who creates a to-do list, assigns tasks to team members, and checks if everything is done.

- **Example:** If you write a PySpark program to count words in a large text file, the Driver translates your code into a plan (e.g., “read file, split words, count them”) and assigns tasks to workers.

2. Cluster Manager

- **What is it?** The Cluster Manager is the resource allocator. It manages the cluster’s resources (CPU, memory) and assigns tasks to worker nodes.
- **Types of Cluster Managers:**
- **Standalone:** Spark’s built-in manager, simple for small clusters.
- **YARN:** Used in Hadoop ecosystems, good for shared clusters.
- **Mesos:** A general-purpose cluster manager.
- **Kubernetes:** Modern option for containerized environments (e.g., on GCP).
- **Key Responsibilities:**
- Allocates resources (e.g., CPU cores, memory) to your Spark application.
- Assigns Executors to worker nodes.
- Monitors resource usage and handles node failures.
- **Analogy:** Think of the Cluster Manager as a factory supervisor who decides which machines (worker nodes) will do the work and ensures they have enough tools (resources).
- **Example:** On Google Cloud’s Dataproc (a Spark cluster service), the Cluster Manager (YARN) assigns Executors to virtual machines in your cluster.

3. Executors

- **What are they?** Executors are worker processes running on worker nodes (computers in the cluster). Each Executor handles a portion of the data and tasks.
- **Key Responsibilities:**
- Execute tasks assigned by the Driver.
- Store data in memory (or disk if memory is full) for fast processing.

- Send results back to the Driver.
- **Analogy:** Executors are like factory workers who follow the manager's instructions, process raw materials (data), and produce results.
- **Example:** If your job is to filter a 1TB dataset, each Executor processes a chunk of the data (e.g., 100GB) in parallel.

4. Worker Nodes

- **What are they?** Physical or virtual machines in the cluster that host Executors.
- **Role:** Provide the computing power (CPU, memory, disk) for Executors to run tasks.
- **Example:** In a GCP Dataproc cluster, worker nodes are Google Compute Engine VMs.

5. SparkContext

- **What is it?** The main entry point for interacting with Spark, created by the Driver.
- **Role:** Manages the connection to the cluster, tracks resources, and coordinates job execution.
- **Example:** In PySpark, you create a `SparkContext` (or `SparkSession` in newer versions) to start your application:

```
from pyspark.sql import SparkSession spark =  
SparkSession.builder.appName("MyApp").getOrCreate()
```

How These Components Work Together

Here's how a Spark job flows through the architecture: 1. You submit a Spark application (e.g., a PySpark script).

2. The **Driver** starts, creates a **SparkContext**, and builds a **DAG** (logical plan) from your code.

3. The Driver communicates with the **Cluster Manager** to request resources.

4. The Cluster Manager allocates **Executors** on **Worker Nodes**.

5. The Driver breaks the DAG into **stages** (groups of tasks) and sends **tasks** to Executors.
6. Executors process tasks in parallel, storing data in memory for speed.
7. Executors send results back to the Driver.
8. The Driver collects results and either outputs them (e.g., to a file) or continues with the next stage.

Key Concepts in Spark Architecture

1. DAG (Directed Acyclic Graph)

- Spark creates a DAG to represent your job as a series of operations (e.g., read, filter, group).
- The DAG is split into **stages**, where each stage contains tasks that can run in parallel.
- **Example:** If you filter a dataset and then group it, Spark creates a DAG with two stages: one for filtering, one for grouping.

2. Stages and Tasks

- A **stage** is a group of tasks that can run without shuffling data across nodes (e.g., filtering rows).
- A **task** is the smallest unit of work, executed by an Executor on a data partition.
- **Example:** If you have 10GB of data split into 10 partitions, Spark creates 10 tasks, each processing 1GB.

3. Data Partitions

- Spark splits large datasets into smaller chunks called **partitions**, processed in parallel by Executors.
- **Example:** A 1TB file might be split into 1000 partitions, each 1GB, processed by multiple Executors.

4. In-Memory Processing

- Spark stores data in memory (RAM) to avoid slow disk reads, making it much faster than Hadoop.

- If memory is full, Spark spills data to disk but optimizes to minimize this.

5. Fault Tolerance

- Spark achieves fault tolerance through **lineage** (the DAG tracks how data was created).
- If an Executor fails, Spark re-runs tasks on another node using the lineage, ensuring no data is lost.
- **Example:** If a node crashes while filtering data, Spark re-executes the filtering task on another node.

• • •

Visualizing Spark's Workflow

Here's a simplified diagram of a Spark job processing a dataset:

Here's a simplified diagram of a Spark job processing a dataset:

```
[Your Code: Read CSV -> Filter -> GroupBy -> Count]
```

```
|
```

```
v
```

```
[Driver: Creates DAG]
```

```
|
```

```
v
```

```
[Cluster Manager: Allocates 3 Executors]
```

```
|
```

```
v
```

```
|
```

```
v
```

```
|
```

```
v
```

```
[Executor 1] [Executor 2] [Executor 3]
```

```
(Process
```

```
(Process
```

```
(Process
```

```
Partition
```

```
Partition
```

```
Partition
```

```
1)
```

```
2)
```

```
3)
```

```
|
```

```
|
```

```
|
```

```
v
```

```
v
```

```
v
```

```
[Results sent back to Driver]
```

```
|
```

```
v
```

```
[Output: Save or Display]
```

Practical Example: Word Count in Spark

Let's see how the architecture works with a classic word count example: — **Goal:** Count the frequency of words in a large text file (e.g., 100GB). — **Code** (in PySpark):

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder.appName("WordCount").getOrCreate()
text = spark.read.text("hdfs://largetextfile.txt")
words = text.rdd.flatMap(lambda line: line[0].split(" "))
wordcounts = words.map(lambda word: (word, 1)).reduceByKey(lambda a, b: a + b)
wordcounts.saveAsTextFile("hdfs://output")
spark.stop()
```

How It Works in the Architecture:

1. Driver:

- Creates a `SparkSession` (includes `SparkContext`).
- Builds a DAG: read file → split into words → count words → save output.
- Splits the DAG into stages (e.g., read + split, count, save).

2. Cluster Manager:

- Allocates, say, 10 Executors on 10 worker nodes.

3. Executors:

- Each Executor processes a partition of the file (e.g., 10GB each).
- Tasks: Split lines into words, count occurrences, aggregate counts.
- Store intermediate data (word counts) in memory.

4. Data Flow:

- Executors send partial counts (e.g., `("hello", 100)`) to the Driver.
- The Driver combines results and saves them to HDFS.

5. Fault Tolerance:

- If an Executor fails, the Driver uses the DAG to re-run tasks on another node.

How a PySpark Job Works When Submitted

Now, let's walk through the end-to-end process of submitting a PySpark job, assuming you're running it on a GCP Dataproc cluster (a common setup for a GCP Data Engineer role).

Step 1: Write the PySpark Code

You write a Python script (`word_count.py`) like the word count example above.

Step 2: Submit the Job

You submit the job to the cluster using the `spark-submit` command:

```
spark-submit --master yarn word_count.py
```

`--master yarn` tells Spark to use YARN as the Cluster Manager (common on Dataproc).

- The script is sent to the Driver node.

Step 3: Driver Initialization

- The **Driver** starts on the master node (or a designated node in the cluster).
- It creates a **SparkSession** (or `SparkContext`) to connect to the cluster.
- The Driver parses your code and builds a **DAG** of operations (e.g., read, transform, save).

Step 4: Resource Allocation

- The Driver contacts the **Cluster Manager** (YARN on Dataproc).
- YARN allocates resources (e.g., 10 Executors with 4GB memory each) across worker nodes.
- Each Executor is a JVM process running on a worker node.

Step 5: DAG and Task Creation

- The Driver divides the DAG into **stages** based on operations that require data shuffling (e.g., `reduceByKey` in word count requires shuffling).
- Each stage is broken into **tasks** (one task per partition).

- The Driver sends tasks to Executors via the Cluster Manager.

Step 6: Task Execution

- Each Executor:
- Receives tasks from the Driver.
- Processes its assigned data partition (e.g., a chunk of the input file).
- Stores intermediate results in memory (e.g., partial word counts).
- Performs transformations (e.g., `flatMap`, `map`) and actions (e.g., `reduceByKey`).
- If a shuffle is needed (e.g., grouping words), Executors exchange data across nodes.

Step 7: Fault Tolerance and Recovery

- If an Executor fails (e.g., node crashes), the Driver detects it via heartbeats.
- The Driver uses the DAG's lineage to re-run failed tasks on another Executor.
- YARN may allocate a new Executor if needed.

Step 8: Result Collection

- Executors send results (e.g., final word counts) back to the Driver.
- The Driver combines results and performs any final actions (e.g., saving to Cloud Storage).

Step 9: Job Completion

- The Driver saves the output (e.g., to `gs://output` on GCP).
- The SparkSession is closed, and resources are released by the Cluster Manager.
- You see the job's output or logs in the Dataproc UI or command line.

Example on GCP Dataproc

- You create a Dataproc cluster with 1 master node and 5 worker nodes.
- Submit the job: `gcloud dataproc jobs submit pyspark --cluster=my-cluster word_count.py`.
- YARN allocates Executors to worker nodes.

- The job processes a 100GB file stored in Google Cloud Storage (`gs://input`).
- Output is saved to `gs://output` .

Summary

- **Spark Architecture:**
- **Driver:** Plans and coordinates the job, creates DAG.
- **Cluster Manager:** Allocates resources (e.g., YARN on Dataproc).
- **Executors:** Process data in parallel on worker nodes.
- **SparkContext:** Connects your program to the cluster.
- **PySpark Job:**
- Submit code → Driver builds DAG → Cluster Manager allocates Executors → Tasks run in parallel → Results collected → Output saved.
- **Key Features:** In-memory processing, fault tolerance via lineage, scalability for big data.

Note

If this article helped you gain some knowledge, please clap and comment. Don't forget to follow me on [Medium](#) . Your support helps me create more content like this and keeps us connected in the data engineering community. Thank you!

Pyspark

Spark Architecture

Apache Spark

Data Engineering



Follow

Written by DataWithSantosh

253 followers · 24 following

A GCP Data Engineer sharing cutting-edge data insights. For the latest in data engineering!